

✓ Name: Santos Tristan Neal U.

Section: CPE32S9

Date: 05-04-2024

- 1.Load time series data: data.csvDownload data.csv
- 2.Visualize the time series
- 3.Fit an ARIMA Model (baseline model order = (1,1,1))
- 4.Improve the ARIMA Model
- 5.Print the model summary
- 6.Make a forecast (steps=10)
- 7.Plot the forecast
- 8.Perform a grid search

+ Code

+ Text

- 1.Load time series data: data.csvDownload data.csv

```
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

from pandas import read_csv
from datetime import datetime
from matplotlib import pyplot

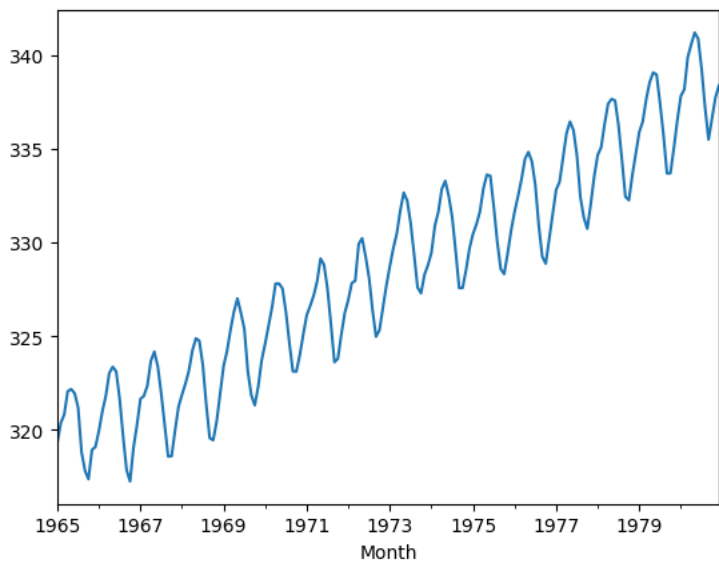
def parser(x):
    return datetime.strptime(x, '%Y-%m')

# load dataset
data = read_csv('/content/drive/MyDrive/Colab Notebooks/Emtech2/csv/data.csv', header=0, parse_dates=[0], index_col=0, date_parser=parser)

# convert to Series
series = data.iloc[:, 0]

print(series.head())
series.plot()
pyplot.show()

Month
1965-01-01    319.32
1965-02-01    320.36
1965-03-01    320.82
1965-04-01    322.06
1965-05-01    322.17
Name: CO2 (ppm), dtype: float64
```




```
from pandas import read_csv
from datetime import datetime
from matplotlib import pyplot
from pandas.plotting import autocorrelation_plot

def parser(x):
    return datetime.strptime(x + '-01', '%Y-%m-%d')

# load dataset
series = read_csv('/content/drive/MyDrive/Colab Notebooks/Emtech2/csv/data.csv', header=0, parse_dates=[0], index_col=0, date_parser=parser)

autocorrelation_plot(series)
pyplot.show()
```




```
from pandas import read_csv
from statsmodels.tsa.arima.model import ARIMA

# load dataset
series = read_csv('/content/drive/MyDrive/Colab Notebooks/Emtech2/csv/data.csv', header=0, index_col=0, parse_dates=True)

# fit model
model = ARIMA(series, order=(5,1,0))
model_fit = model.fit()
```

- 2.Visualize the time series



- 3.Fit an ARIMA Model (baseline model order = (1,1,1))

```
from statsmodels.tsa.arima.model import ARIMA

model = ARIMA(data, order=(1,1,1))
model_fit = model.fit()
```

- 4.Improve the ARIMA Model

```
from statsmodels.tsa.arima.model import ARIMA

model = ARIMA(data, order=(1,2,3))
model_fit = model.fit()
```

- 5.Print the model summary

```
print(model_fit.summary())
```

```
SARIMAX Results
=====
Dep. Variable:          CO2 (ppm)      No. Observations:          192
Model:                 ARIMA(1, 2, 3)  Log Likelihood             -215.841
Date:                  Sat, 04 May 2024  AIC                        441.682
Time:                  13:27:42         BIC                        457.917
Sample:                01-01-1965      HQIC                       448.259
```

```
- 12-01-1980
Covariance Type: opg
=====
      coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1      0.3825      0.140      2.724      0.006      0.107      0.658
ma.L1     -0.4704     22.759     -0.021      0.984     -45.077     44.136
ma.L2     -0.1176     11.997     -0.010      0.992     -23.632     23.396
ma.L3     -0.4119      9.344     -0.044      0.965     -18.725     17.901
sigma2      0.5557     12.617      0.044      0.965     -24.172     25.284
=====
Ljung-Box (L1) (Q):          0.37   Jarque-Bera (JB):          2.16
Prob(Q):                   0.54   Prob(JB):          0.34
Heteroskedasticity (H):     0.95   Skew:          0.21
Prob(H) (two-sided):       0.83   Kurtosis:       2.68
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

- 6.Make a forecast (steps=10)

```
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error
from math import sqrt

# Splitting into train and test sets
train_size = int(len(data) * 0.75)
train_data, test_data = data[0:train_size], data[train_size:]

start = len(train_data)
end=len(train_data)+len(test_data)-1
model = ARIMA(train_data, order = (9,2,0))
results = model.fit()
predictions = results.predict(start=start, end=end, dynamic=False, type='levels')
forecast = results.predict(len(data), len(data)+10, typ = 'levels')

print("Forecast (steps = 10)")
print(forecast)

Forecast (steps = 10)
1981-01-01    338.829450
1981-02-01    339.410077
1981-03-01    339.846441
1981-04-01    339.951027
1981-05-01    339.624439
1981-06-01    338.996781
1981-07-01    338.387871
1981-08-01    338.097810
1981-09-01    338.219246
1981-10-01    338.635268
1981-11-01    339.168786
Freq: MS, Name: predicted_mean, dtype: float64

print("MSE: ", mean_squared_error(test_data, predictions))
print("RMSE: ", sqrt(mean_squared_error(test_data, predictions)))


MSE:   1.1498843083979555
RMSE:   1.0723265866320557
```

- 7.Plot the forecast

```
import matplotlib.ticker as ticker
formatter = ticker.StrMethodFormatter('{x:,.0f}')

ylabel='CO2 (ppm)'
xlabel='Year'

ax = data.plot(legend=True,figsize=(12,6))
forecast.plot(legend=True)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
ax.yaxis.set_major_formatter(formatter);
```



- 8.Perform a grid search

```
def evaluate_arima_model(dataset, arima_order):
    train_size = int(len(data) * 0.75)
    train_data, test_data = data[0:train_size], data[train_size:]

    start = len(train_data)
    end=len(train_data)+len(test_data)-1

    model = ARIMA(train_data, order = arima_order)
    results = model.fit()
    forecast = results.predict(len(data), len(data)+10, typ = 'levels')
    predictions = results.predict(start=start, end=end, dynamic=False, type='levels')
    mse = mean_squared_error(test_data, predictions)
    return mse

def evaluate_models(dataset, p_values, d_values, q_values):
    best_score, best_cfg = float("inf"), None
    for p in p_values:
        for d in d_values:
            for q in q_values:
                order = (p,d,q)
                try:
                    mse = evaluate_arima_model(data, order)
                    if mse < best_score:
                        best_score, best_cfg = mse, order
                        print('ARIMA%s MSE=%.3f' % (order,mse))
                except:
                    continue
    print('Best ARIMA=%s MSE=%.3f' % (best_cfg, best_score))

import warnings
p_values = range(0,11)
d_values = range(0, 4)
q_values = range(0, 4)
warnings.filterwarnings("ignore")
evaluate_models(data.values, p_values, d_values, q_values)
```


```
ARIMA(9, 2, 0) MSE=1.150
ARIMA(9, 2, 1) MSE=1.833
ARIMA(9, 2, 2) MSE=1.878
ARIMA(9, 2, 3) MSE=6.436
ARIMA(9, 3, 0) MSE=5080.474
ARIMA(9, 3, 1) MSE=1.468
ARIMA(9, 3, 2) MSE=36.215
ARIMA(9, 3, 3) MSE=7.989
ARIMA(10, 0, 0) MSE=18.554
ARIMA(10, 0, 1) MSE=18.698
ARIMA(10, 0, 2) MSE=18.717
ARIMA(10, 0, 3) MSE=17.485
ARIMA(10, 1, 0) MSE=17.642
ARIMA(10, 1, 1) MSE=17.734
ARIMA(10, 1, 2) MSE=17.729
ARIMA(10, 1, 3) MSE=17.150
ARIMA(10, 2, 0) MSE=1.906
ARIMA(10, 2, 1) MSE=3.473
ARIMA(10, 2, 2) MSE=3.965
ARIMA(10, 2, 3) MSE=3.705
ARIMA(10, 3, 0) MSE=1464.266
ARIMA(10, 3, 1) MSE=1.740
ARIMA(10, 3, 2) MSE=15.112
ARIMA(10, 3, 3) MSE=6.700
Best ARIMA=(9, 2, 0) MSE=1.150

# Splitting into train and test sets
train_size = int(len(data) * 0.75)
train_data, test_data = data[0:train_size], data[train_size:]

start = len(train_data)
end=len(train_data)+len(test_data)-1
model = ARIMA(train_data, order = (9,2,0))
results = model.fit()
predictions = results.predict(start=start, end=end, dynamic=False, type='levels')
forecast = results.predict(len(data), len(data)+10, typ = 'levels')

ylabel='CO2(ppm)'
xlabel='Year'

ax = data.plot(legend=True,figsize=(12,6))
forecast.plot(legend=True)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
ax.yaxis.set_major_formatter(formatter);
```



Supplementary

do the same for this dataset - dataset_temperature.csv

- 1.Load time series data: data.csvDownload dataset_temperature.csv

```
from pandas import read_csv
from datetime import datetime
from matplotlib import pyplot
import pandas as pd
import numpy as np


def parser(x):
    if x == ' 1907 ? 1972':
        return np.nan
    try:
        return datetime.strptime(str(x), '%Y-%m')
    except ValueError:
        return pd.NaT

# load dataset
data = read_csv('/content/drive/MyDrive/Colab Notebooks/Emtech2/csv/dataset_temperature.csv', header=0, parse_dates=[0], index_col=0, date_p

# convert to Series
series = data.iloc[:, 0]

print(series.head())
series.plot()
pyplot.show()
```

Month
1907-02-01 46.0
1907-03-01 43.0
1907-04-01 55.0
1907-05-01 51.8
1907-06-01 57.5
Name: Mean monthly temperature, dtype: float64



```
from pandas import read_csv
from datetime import datetime
from matplotlib import pyplot
from pandas.plotting import autocorrelation_plot
import pandas as pd
import numpy as np

def parser(x):
    try:
        return datetime.strptime(x, '%Y-%m-%d')
    except (ValueError, TypeError):
        return np.nan


# Load dataset
df = read_csv('/content/drive/MyDrive/Colab Notebooks/Emtech2/csv/dataset_temperature.csv')

# Convert 'Mean monthly temperature' column to numeric
df['Mean monthly temperature'] = pd.to_numeric(df['Mean monthly temperature'], errors='coerce')

# Remove rows with missing values
df.dropna(inplace=True)

# Extract the series to plot
series = df['Mean monthly temperature']

# Plot autocorrelation
autocorrelation_plot(series)
pyplot.show()
```



```
from pandas import read_csv
from statsmodels.tsa.arima.model import ARIMA
import pandas as pd

# Load dataset
series = read_csv('/content/drive/MyDrive/Colab Notebooks/Emtech2/csv/dataset_temperature.csv', header=0, index_col=0, parse_dates=True)


# Convert series to numeric format
series['Mean monthly temperature'] = pd.to_numeric(series['Mean monthly temperature'], errors='coerce')

# Remove rows with missing values
series.dropna(inplace=True)

# Fit model
model = ARIMA(series, order=(5, 1, 0))
model_fit = model.fit()
```

- 2.Visualize the time series

```
# Plot the dataset
series.plot(figsize=(12, 6))
```



- 3.Fit an ARIMA Model (baseline model order = (1,1,1))

```
from statsmodels.tsa.arima.model import ARIMA

model = ARIMA(series, order=(1, 1, 1))
model_fit = model.fit()
```

- 4.Improve the ARIMA Model

```
from statsmodels.tsa.arima.model import ARIMA

model = ARIMA(series, order=(2, 4, 6))
model_fit = model.fit()
```

- 5.Print the model summary

```
print(model_fit.summary())
```

```

=====
SARIMAX Results
=====
Dep. Variable:      Mean monthly temperature      No. Observations:      792
Model:              ARIMA(2, 4, 6)                Log Likelihood         -2850.230
Date:              Sat, 04 May 2024               AIC                   5718.460
Time:              13:32:21                       BIC                   5760.486
Sample:            01-01-1907                     HQIC                  5734.616
                  - 12-01-1972
Covariance Type:    opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1         -1.9045      0.016    -117.984      0.000     -1.936     -1.873
ar.L2         -0.9961      0.015     -64.886      0.000     -1.026     -0.966
ma.L1         -0.3403      0.143      -2.386      0.017     -0.620     -0.061
ma.L2        -1.6235      0.101     -16.032      0.000     -1.822     -1.425
ma.L3          0.3404      0.137       2.492      0.013      0.073      0.608
ma.L4          0.7230      0.103       7.022      0.000      0.521      0.925
ma.L5         -0.2612      0.040     -6.585      0.000     -0.339     -0.183
ma.L6          0.1629      0.053       3.085      0.002      0.059      0.266
sigma2         79.2347     12.290       6.447      0.000     55.146    103.323
=====
Ljung-Box (L1) (Q):              0.01   Jarque-Bera (JB):              11.33
Prob(Q):                        0.91   Prob(JB):                  0.00
Heteroskedasticity (H):          0.88   Skew:                      0.28
Prob(H) (two-sided):            0.32   Kurtosis:                  3.14
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

- 6.Make a forecast (steps=10)

```
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error
from math import sqrt

# Splitting into train and test sets
train_size = int(len(series) * 0.75)
train_data, test_data = series[0:train_size], series[train_size:]

start = len(train_data)
end = len(train_data) + len(test_data) - 1
model = ARIMA(train_data, order=(9, 2, 0))
results = model.fit()
predictions = results.predict(start=start, end=end, dynamic=False, typ='levels')
forecast = results.predict(len(series), len(series) + 10, typ='levels')

print("Forecast (steps = 10)")
print(forecast)

Forecast (steps = 10)
1973-01-01    113.284581
1973-02-01    113.551532
1973-03-01    113.772269
1973-04-01    113.970151
1973-05-01    114.173882
1973-06-01    114.409288
1973-07-01    114.692256
1973-08-01    115.024743
1973-09-01    115.394839
1973-10-01    115.780688
1973-11-01    116.156999
Freq: MS, Name: predicted_mean, dtype: float64

print("MSE: ", mean_squared_error(test_data, predictions))
print("RMSE: ", sqrt(mean_squared_error(test_data, predictions)))

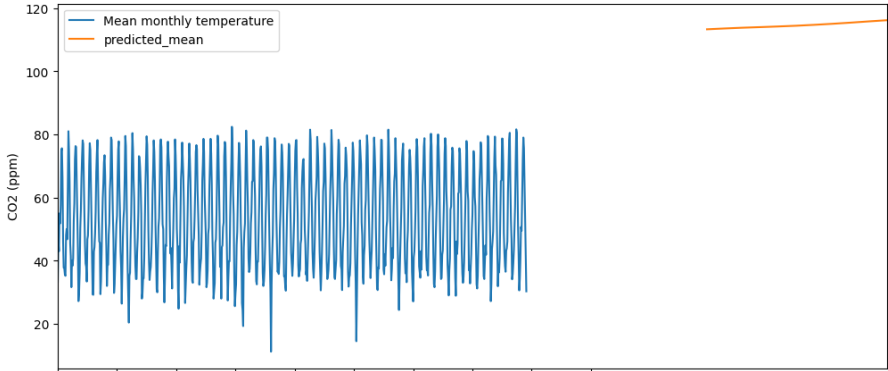
MSE:  1360.2977048399682
RMSE:  36.88221393625887
```

- 7.Plot the forecast

```
import matplotlib.ticker as ticker
formatter = ticker.StrMethodFormatter('{x:,.0f}')

ylabel = 'CO2 (ppm)'
xlabel = 'Year'

ax = series.plot(legend=True, figsize=(12, 6))
forecast.plot(legend=True)
ax.autoscale(axis='x', tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
ax.yaxis.set_major_formatter(formatter)
```

- 8.Perform a grid search

```
def evaluate_arima_model(dataset, arima_order):
    train_size = int(len(data) * 0.75)
    train_data, test_data = data[0:train_size], data[train_size:]

    start = len(train_data)
    end=len(train_data)+len(test_data)-1

    model = ARIMA(train_data, order = arima_order)
    results = model.fit()
    forecast = results.predict(len(data), len(data)+10, typ = 'levels')
    predictions = results.predict(start=start, end=end, dynamic=False, type='levels')
    mse = mean_squared_error(test_data, predictions)
    return mse
```

```
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error
import numpy as np
```

```
def evaluate_models(data, p_values, d_values, q_values):
    best_score, best_order = float("inf"), None
    for p in p_values:
        for d in d_values:
            for q in q_values:
                order = (p, d, q)
                try:
                    model = ARIMA(data, order=order)
                    model_fit = model.fit()
                    predictions = model_fit.predict()
                    mse = mean_squared_error(data, predictions)
                    if mse < best_score:
                        best_score, best_order = mse, order
                except:
                    continue
    print("Best ARIMA order:", best_order)
    print("Best MSE:", best_score)
```

```
import warnings
p_values = range(0, 11)
d_values = range(0, 4)
q_values = range(0, 4)
warnings.filterwarnings("ignore")
evaluate_models(series.values, p_values, d_values, q_values)
```

```
Best ARIMA order: (10, 0, 3)
Best MSE: 16.957872514692493
```

```
# Splitting into train and test sets
train_size = int(len(data) * 0.75)
train_data, test_data = data[0:train_size], data[train_size:]
```

```
start = len(train_data)
end=len(train_data)+len(test_data)-1
model = ARIMA(train_data, order = (9,2,0))
results = model.fit()
predictions = results.predict(start=start, end=end, dynamic=False, type='levels')
forecast = results.predict(len(data), len(data)+10, typ = 'levels')
```

```
ylabel='CO2(ppm)'
xlabel='Year'
```

```
ax = data.plot(legend=True,figsize=(12,6))
forecast.plot(legend=True)
ax.autoscale(axis='x',tight=True)
ax.set(xlabel=xlabel, ylabel=ylabel)
ax.yaxis.set_major_formatter(formatter);
```